

Decision Log : Drone Operations AI

Key Assumptions

Data & Storage

CSV-based persistence - Assumed data is read at startup, not modified during runtime

- Valid for prototype with 5-10 records
- No concurrent write requirements
- Trade-off: Simplicity over scalability

Single-threaded operation - One user/conversation at a time

- Sufficient for demo/proof-of-concept
- Can upgrade to async with FastAPI + Celery later

Pilot/drone status is current - Data manually updated in CSV before queries

- No real-time location tracking or live updates
- Assumption: Operators trust CSV data accuracy

User Interaction

Natural language queries require LLM intelligence

- Fallback: Simple regex pattern matching for ~20 common queries
- Assumption: Users phrase requests consistently ("Show pilots", "Available drones")
- Reality check: Rule-based mode covers 80% of expected operational queries

No authentication needed - Suitable for prototype/demo only

- Assumption: Internal tool for trusted operators
- Future: Add JWT tokens for production

Operational Model

Urgent missions detected but require human approval (NOT auto-reassignment)

- Assumption: Safety-critical drone operations require human oversight
- Liability: Auto-assigning incompetent pilot = legal/regulatory risk
- Model: System detects conflicts and alerts operator

Technology Trade-offs & Decisions

1. Flask vs FastAPI vs Django

Framework	Chosen?	Reason
Flask	YES	Lightweight, perfect for prototype-to-production, single file deployable, minimal overhead
FastAPI	NO	Overkill for simple API, steeper learning curve, async not needed yet
Django	NO	Too heavy, 20+ files boilerplate, better for complex apps with admin panels

Trade-off Accepted: Limited built-in features, manual error handling
Benefit: Flexible, easy to scale to production (can migrate to FastAPI/Django later)

2. LLM: ChatGPT + Rule-Based Fallback

Approach	Chosen?	Reason
ChatGPT + Fallback	YES	Best of both: intelligent LLM + graceful degradation when API fails
Claude Only	NO	More reliable but no fallback, requires API key
Local LLM (Llama)	NO	Free & offline but GPU-intensive, slower, overkill for structured queries
Rule-Based Only	NO (as primary)	Simple but limited, can't handle context or variations

Dual-Mode Architecture:

```
User Query → [Is OpenAI API available?]
└ YES → ChatGPT with Tool descriptions (powerful reasoning)
└ NO → Regex pattern matching (simple, fast, 100% reliable)
```

Why This Works:

Real scenario: OpenAI API quota exceeded (actual current state) → App still works perfectly
LLM mode: Better NLP, conversation context, reasoning
Rule-based mode: 80% of operational queries work without AI
Trade-off: Rule-based is regex-fragile but sufficient for structured commands

Storage	Chosen?	Reason
CSV	YES	Human-readable, version control friendly, zero setup, lightweight
SQLite	NO (for now)	Overkill for 12 records, still needs schema management
PostgreSQL	NO (for now)	Needs docker/server, connection pooling, overcomplicated for MVP
Google Sheets	NO (Phase 2)	Adds OAuth complexity, but planned for real 2-way sync future

3. Data Storage: CSV vs SQLite vs PostgreSQL

CSV Limitations Accepted:

No concurrent write safety
No indexing (but 5 pilots = linear search fine)
No complex queries

But: Simple, git-trackable, editable in Excel

Phase Plan:

- Phase 1: CSV (current)
- Phase 2: SQLite + ORM
- Phase 3: PostgreSQL with replicas

4. Conflict Detection Strategy

Chosen: Validation-based (proactive) + user approval required

5 Conflict Types Detected:

- Double-Booking** - Same pilot assigned to overlapping missions
- Skill Mismatch** - Pilot lacks required skills
- Certification Gap** - Missing regulatory certifications
- Equipment Mismatch** - Drone lacks required capabilities
- Location Mismatch** - Resource unavailable at mission location

Why Validation-Based (not auto-resolve):

- Safety-first:** Prevents assigning unqualified pilots
- Regulatory:** Drone ops require documented human decision
- Liability:** "System auto-assigned incompetent pilot" = lawsuit

How "Urgent Reassignments" Were Interpreted

What We Did NOT Do

- Auto-reassign pilots without approval
- Override conflict detection for urgent missions
- Implement automatic conflict resolution

What We DID Do

- Detect urgent missions (priority field)
- Flag conflicts immediately when urgent mission is affected
- Return actionable suggestions: "Reassign pilot X to mission Y to free up Z"
- Require human dispatcher to make final decision

Why Monolithic App (Not Microservices)?

- Current:** Single Flask app with all logic
- Benefits:** Single docker deployment, no inter-service latency, easy to debug
- Limitation:** Can't scale conflict detection independently
- Scale Plan:** Split into Agent Service + Conflict Service + DB Service when needed

Why Rule-Based Fallback > Hard Failure?

- Alternative: Fail with error if API down
- Our choice: Gracefully degrade to pattern matching
- Real-world impact: Demo at 2am when OpenAI quota exceeded? Works anyway ✓
- User experience: No API key requirement for basic queries